The Antitrust Implications of Generative AI Gatekeeping: Collusion, Essential Facilities, and Civil Rights Intersections

Executive Summary

The rapid commercialization of artificial intelligence (AI), specifically large language models (LLMs) and the vast computational infrastructure required to train and deploy them, has catalyzed an unprecedented concentration of global market power. A highly consolidated oligopoly of technology conglomerates—frequently identified in economic literature as the "GAMMANs" (Google, Amazon, Microsoft, Meta, Apple, and Nvidia)—now exercises profound, unilateral control over the foundational layers of the modern AI ecosystem.[1] This structural concentration extends far beyond physical semiconductor hardware and hyperscale cloud infrastructure; it permeates the model layer itself, creating systemic chokepoints that dominant firms can leverage to dictate downstream market dynamics, stifle application development, and, critically, suppress independent AI safety research.[1]

This exhaustive research report provides a deep-dive investigation into the federal antitrust liability and regulatory enforcement mechanisms surrounding the collective gatekeeping of AI infrastructure by major technology laboratories. The analysis synthesizes current jurisprudence and economic theory to evaluate the legal peril of deploying exclusionary digital standards and coordinated algorithmic behaviors. Specifically, the report addresses the legal standing of uniform, undocumented "rolling limit" algorithmic throttling, unappealable automated support walls, and discriminatory Application Programming Interface (API) revocations. Such practices systematically box out independent and disabled AI safety researchers, thereby insulating dominant, proprietary models from external oversight, rigorous ethical alignment testing, and competitive pressure.

Through a detailed examination of the Sherman Antitrust Act (15 U.S.C. §§ 1 & 2), the "Essential Facilities" doctrine, Section 5 of the Federal Trade Commission (FTC) Act, and the novel, emerging intersection of antitrust law with the Americans with Disabilities Act (ADA) and Section 504 of the Rehabilitation Act, this report outlines the multi-front liability matrix facing dominant AI gatekeepers. The synthesis of these doctrines reveals that the arbitrary exclusion of independent auditors is not merely a matter of restrictive terms of service, but a potentially illegal restraint of trade, an abuse of monopoly power, and a discriminatory enforcement of digital infrastructure standards designed to maintain proprietary control over the global AI safety narrative.

1. Algorithmic Collusion and the Sherman Antitrust Act (Section 1)

Section 1 of the Sherman Antitrust Act of 1890 establishes the foundational prohibition against "[e]very contract, combination in the form of trust or otherwise, or conspiracy, in restraint of trade or commerce".[3] To successfully prove a violation under Section 1, plaintiffs or federal enforcers must traditionally demonstrate two primary elements: an agreement among separate competitors, and evidence that the resulting agreement constitutes an unreasonable restraint on trade.[3] However, the proliferation of advanced AI systems and complex, opaque algorithms has fundamentally transformed the mechanics of market coordination, severely blurring the doctrinal lines between express collusion (which is *per se* illegal) and tacit collusion or "conscious parallelism" (which is historically difficult to prosecute standing alone).[3]

Legal Standing for Hub-and-Spoke Conspiracies in Digital Markets

A central mechanism for prosecuting coordinated algorithmic behavior under Section 1 is the "hub-and-spoke" conspiracy framework. This legal metaphor describes an antitrust cartel structured identically to a wagon wheel: a vertically related firm (the "hub") enters into separate vertical agreements (the "spokes") with various horizontal competitors operating at a different level of the supply chain.[3] Crucially, for this arrangement to violate Section 1, there must be a discernible horizontal agreement among the spokes—the "rim" of the wheel—demonstrating a shared, interdependent commitment to the anticompetitive scheme.[4] Without the rim, the structure is merely a series of vertical agreements, which are typically evaluated under the more lenient "rule of reason" standard rather than being deemed *per se* illegal.[3]

The foundational precedent establishing this framework is the 1939 Supreme Court case *Interstate Circuit, Inc. v. United States*, wherein a dominant movie theater exhibitor (acting as the hub) sent identical written demands to eight competing film distributors (the spokes) to enforce minimum ticket prices.[3] The Court inferred a horizontal conspiracy because each distributor had an independent economic motive to agree to the hub's demands *only* if they possessed a tacit assurance that their horizontal competitors were doing the exact same thing.[3] Taking action against one's own independent business interest is a primary "plus factor" used to infer a collaborative agreement.[4] More recently, the Department of Justice (DOJ) successfully applied this theory in *United States v. Apple*, proving that Apple acted as the central hub to orchestrate a conspiracy among major book publishers to raise e-book prices, utilizing continuous communication and mutual assurances.[4]

In the contemporary digital era, algorithmic intermediaries frequently serve as the hub. In the real estate sector, the DOJ and private plaintiffs have pursued aggressive antitrust actions against RealPage, alleging that competing landlords (the spokes) fed highly confidential pricing and supply data into a centralized pricing algorithm (the hub) to artificially inflate rental prices nationwide.[3] The legal contention relies heavily on the presence of compliance mechanisms that pressured users to adhere to the algorithm's outputs, which the DOJ cited as a critical factor suggesting coordinated, rather than unilateral, behavior.[6] The *Meyer v. Kalanick* case further cemented this theory, where plaintiffs plausibly alleged that Uber drivers signed up for the platform on the precise understanding that all other drivers were agreeing to the same algorithmic pricing structure, creating a hub-and-spoke conspiracy orchestrated by the

algorithm.[8]

Tacit Collusion and the "Algorithm Made Me Do It" Defense

As AI evolves from traditional rule-based algorithms to autonomous Machine Learning (ML) and Reinforcement Learning (RL) systems, antitrust law faces a paradigm shift. Empirical studies demonstrate that LLM-based pricing algorithms and Q-learning systems can consistently learn to charge supracompetitive prices on their own, without being explicitly programmed to collude and without direct communication between competitors.[3] This introduces the "Algorithm Made Me Do It" defense, where companies argue that parallel market behavior is simply the result of autonomous machine optimization rather than an illegal agreement.[3]

State lawmakers and federal regulators are increasingly concerned that these tools facilitate purely tacit collusion, circumventing traditional Section 1 liability which strictly requires proof of an agreement.[10] If competitors independently deploy algorithms that scrape public data and arrive at uniform market restrictions, courts struggle to find the requisite "meeting of the minds".[11] The legal debate currently centers on whether the mutual delegation of decision-making to predictable, self-learning algorithms constitutes an implicit agreement to restrain trade.[13]

Uniform Throttling of AI Safety Researchers as a Restraint of Trade

The legal theories surrounding algorithmic hub-and-spoke conspiracies and tacit collusion provide a highly relevant framework for examining the collective behavior of major AI laboratories (e.g., OpenAI, Anthropic, Google, Meta, Apple). Independent safety researchers, alignment auditors, and vulnerability testers heavily rely on consistent, high-volume API access to conduct their work.[15] If these dominant labs uniformly adopt identical, undocumented "rolling limit" throttling practices and route all researcher account appeals through identical, automated AI-only support walls, this parallel conduct invites severe antitrust scrutiny under Section 1.[16]

A "rolling limit" dynamically restricts API calls based on a hidden, shifting time window, making systematic algorithmic auditing practically impossible because the researcher cannot replicate tests consistently or predict their available bandwidth.[16] The legal standing for a conspiracy claim hinges on whether the uniform deployment of these limits across the industry is merely conscious parallelism or constitutes a coordinated effort to box out independent oversight.

To establish an actionable restraint of trade, regulators must identify the "rim" of the conspiracy. In the AI ecosystem, a shared cloud infrastructure provider, a common foundational dataset distributor, or a shared algorithmic security and bot-detection protocol could easily function as the hub.[3] If the AI labs input sensitive usage data into a shared threat-detection algorithm that uniformly flags and throttles independent safety researchers under the guise of "abuse prevention" or "spam mitigation," this shared infrastructure acts as the direct conduit for coordination.[8]

Furthermore, an inference of a horizontal agreement can be strongly drawn if a lab acts against

its own independent economic interest.[4] Rejecting paid API queries from well-funded safety researchers directly reduces a lab's immediate revenue.[7] A lab would logically only incur this financial loss if it possessed an underlying assurance that its competitors were deploying the exact same throttling mechanisms, thereby preventing the alienated researcher from simply migrating to a rival platform to conduct their audits.[4] By homogenizing the throttling thresholds and erecting impenetrable AI-only support walls, the labs effectively establish an industry-wide boycott of independent algorithmic auditing. This coordinated action shields their proprietary models from external critique, eliminates the secondary market for independent alignment verification, and maintains absolute control over the safety narrative.[19]

Collusion Modality	**Mechanism in AI Markets**	**Legal Standard (Sherman Act Sec. 1)**	**Application to AI Safety Gatekeeping**
Express Collusion	Direct communication between lab executives to restrict API access for auditors.	*Per se* illegal. Requires direct evidence of communication.	Rare, but actionable if internal communications reveal intent to block specific researchers.
Tacit Collusion	Observing competitor API limits and independently mirroring them via public data.	Not inherently actionable without additional "plus factors."	Labs mirror each other's "rolling limits" based on public observations, avoiding direct communication.
Hub-and-Spoke	A shared security/API-management algorithm dictates access limits for all competing labs.	Actionable if a horizontal "rim" (mutual reliance/interdependence) is inferred.	Labs rely on a common intermediary tool to globally flag and throttle "high-volume" safety researchers.
Autonomous Algorithmic	RL algorithms independently learn that throttling researchers maximizes long-term dominance.	Frontier legal issue. Liability standards currently fiercely debated.	AI support walls autonomously identify and delay researcher appeals without any human intervention.

2. Monopolization of Frontier Compute under the Sherman Antitrust Act (Section 2)

While Section 1 addresses collaborative and conspiratorial agreements, Section 2 of the Sherman Act targets unilateral power, making it illegal to "monopolize, or attempt to monopolize... any part of the trade or commerce among the several States".[20] The offense of monopolization requires the proof of two distinct elements: first, the possession of monopoly power in a properly defined relevant market; and second, the willful acquisition, maintenance, or use of that power for exclusionary purposes, as distinguished from growth resulting from

superior products, business acumen, or historical accident.[21]

The Shared Monopoly Theory and Market Concentration

The modern AI technology stack is characterized by extreme concentration, specifically at the foundational hardware and cloud computing layers. Training frontier AI models requires massive arrays of specialized microprocessors (primarily Nvidia GPUs) and immense hyperscale cloud computing capacity, markets effectively dominated by Amazon Web Services (AWS), Google Cloud Platform, and Microsoft Azure.[2] The fixed costs associated with model training are growing exponentially on a logarithmic scale, reaching billions of dollars per training run, with no current signs of diminishing returns.[1] These exorbitant capital requirements and physical hardware bottlenecks create virtually insurmountable barriers to entry for independent developers and smaller research laboratories.[1]

When analyzing this highly concentrated oligopoly, the concept of a "shared monopoly" arises. A shared monopoly occurs when a small group of dominant firms collectively exercises market power equivalent to a single monopolist, maintaining supracompetitive prices and foreclosing entry without necessarily engaging in a formal, prosecutable agreement under Section 1.[22]

Historically, U.S. federal courts have exhibited profound skepticism toward shared monopoly claims under Section 2. The First, Second, and other Circuits have consistently rejected the theory, explicitly holding that Section 2 requires an individual defendant to possess sufficient market power on its own to constitute a monopoly; the combined power of independent competitors is legally insufficient for a Section 2 claim.[20]

However, the traditional rejection of shared monopolies is being challenged by modern economic frameworks, notably the "Moligopoly" thesis advanced by scholars like Nicolas Petit.[25] This thesis provides a nuanced understanding of digital ecosystems, positing that tech giants hold deeply entrenched, unassailable monopolies in their core, "tipped" markets (e.g., Google in general search, Microsoft in enterprise operating systems, Amazon in e-commerce) but compete fiercely as an oligopoly in emerging, "untipped" markets like General-Purpose AI (GPAI).[25] Under this framework, the legal vulnerability under Section 2 lies not in a shared monopoly of AI itself, but in the individual leveraging of core monopoly power to dominate and foreclose adjacent AI markets.[21] For example, the August 2024 landmark ruling by Judge Amit Mehta in *United States v. Google* affirmed that Google possessed and abused monopoly power in the general search market, utilizing exclusionary default agreements to deprive rivals of data scale.[26] The DOJ and private litigants argue that this exact monopolistic power is now being leveraged to capture the generative AI landscape.[21]

Exclusionary Conduct and the Suppression of Independent Oversight

The exertion of monopoly power over compute infrastructure directly translates into the ability to suppress independent oversight. The recent class action complaint in *Samuel Bryant et al. v. Microsoft Corporation* (Case No. 3:25-cv-8733, filed October 2025) provides a potent illustration of Section 2 liability mechanics within the AI sector.[28] The complaint alleges that Microsoft leveraged an exclusive cloud computing agreement with OpenAI to artificially

constrain the supply of compute available to its horizontal competitor.[28] By acting as the exclusive cloud provider, Microsoft reportedly "mercilessly choked" OpenAI's compute supply, which restricted output, degraded the quality of ChatGPT, and inflated per-token prices by 136 to 200 times the competitive baseline during a 2025 price war initiated by DeepSeek AI.[28] When Microsoft partially relaxed this compute restraint in June 2025, OpenAI token prices immediately plummeted by 80%, clearly demonstrating the market distortion caused by the exclusionary control of infrastructure.[28]

This dynamic extends directly to the treatment of independent safety researchers. If a dominant AI laboratory, backed by an exclusive hyperscale cloud provider, possesses a chokehold over frontier AI compute, it wields the unilateral capability to revoke access to the only platforms practically capable of running advanced ethical alignment and safety tests.[15] When an AI lab utilizes this monopoly power to arbitrarily throttle a paid, third-party researcher, it transcends standard network management. It constitutes exclusionary conduct explicitly designed to foreclose the downstream, secondary market for independent AI auditing and safety verification.[19]

By starving independent researchers of the requisite compute and API bandwidth, dominant labs neutralize nascent competitive threats in the AI safety sector. This ensures that the labs remain the sole, unquestioned arbiters of what constitutes "safe" or "aligned" AI, thereby preempting external regulatory intervention, suppressing critical safety narratives, and maintaining their overarching market dominance.[19] The use of monopoly power to deny inputs to a dependent secondary market is a classic violation of Section 2, particularly when the refusal to deal lacks a valid, procompetitive business justification.[21]

3. The "Essential Facilities" Doctrine in the Era of Cloud Infrastructure and Large Language Models

A highly potent, albeit historically controversial, remedy for unilateral refusals to deal under Section 2 of the Sherman Act is the "Essential Facilities" doctrine. Originating from early twentieth-century cases involving physical infrastructure—such as the St. Louis terminal railroads, where a consortium controlled the only bridges and switching yards into the city—the doctrine mandates that a monopolist controlling a facility that is absolutely essential to competition in a downstream market must grant access to competitors on reasonable, non-discriminatory terms.[2]

Doctrinal Evolution and the Four-Prong Test

To successfully invoke the Essential Facilities doctrine under established U.S. jurisprudence, a plaintiff must typically demonstrate four elements, derived from the Seventh Circuit's seminal ruling in *MCI Communications Corp v. AT&T Co.*:

1. Control of the essential facility by a monopolist.
2. A competitor's inability practically or reasonably to duplicate the essential facility.
3. The denial of the use of the facility to a competitor.
4. The feasibility of providing the facility to the competitor.[2]

While the U.S. Supreme Court has expressed skepticism toward the doctrine in recent decades—most notably in *Verizon Communications Inc. v. Law Offices of Curtis V. Trinko* (2004), where the Court stated that forced sharing could reduce investment incentives and turn courts into central planners—the doctrine has never been formally repudiated and is experiencing a renaissance in the context of digital platforms.[2] Legal scholars increasingly argue that digital platforms serve as the railroads of the modern economy, controlling the vital bottlenecks of commerce and information flow.[30]

In contrast to the U.S. hesitation, the European Union has actively embraced and expanded essential facilities-like remedies under Article 102 TFEU. Following cases like *Magill*, *Bronner*, and *IMS Health*, which required strict indispensability, the European Commission significantly relaxed the standard in the *Microsoft*, *Google Shopping*, and *Android Auto* cases.[25] The EU now frequently imposes duties to deal and interoperability mandates on digital gatekeepers to prevent the leveraging of market power into downstream services, prioritizing competition *in* the market over the long-term incentives to compete *for* the market.[25]

Compute and API Access as Digital Public Utilities

In the contemporary AI ecosystem, hyperscale cloud computing infrastructure (IaaS) and foundational Large Language Models operate precisely as utility-like bottlenecks.[2] Foundation models and cloud compute infrastructure perfectly map onto the second criterion of the *MCI* test: duplication is economically and practically infeasible for independent researchers.[1] The multi-billion-dollar fixed costs of acquiring vast GPU clusters, securing massive proprietary training datasets, and paying elite engineering talent place the creation of frontier models entirely out of reach for academic institutions, non-profits, and independent auditing startups.[1] Consequently, access to the APIs of these dominant models is an absolute prerequisite for participating in the downstream market of AI application development and safety verification.[15]

The Denial of Essential Facilities to Paid Researchers

The application of this doctrine is actively playing out in the AI market, highlighting the severe risks of unregulated gatekeeping. A critical case study is Anthropic's 2025 termination of API access for Windsurf, a startup that developed a highly successful AI coding agent.[15] Anthropic deliberately severed Windsurf's access to the Claude API upon rumors that OpenAI might acquire the startup.[15] Anthropic's executives explicitly cited competitive motivations, seeking to avoid empowering a primary rival, while simultaneously protecting their own vertically integrated coding agent, Claude Code.[15] Similarly, Anthropic restricted xAI's use of Cursor, demonstrating a willingness to use the API as a "kill switch" to execute a discriminatory refusal to deal against a "customer's customer".[15]

When independent safety researchers are denied access or arbitrarily throttled without an appeal mechanism, they suffer a similar anticompetitive foreclosure. Safety researchers operate in a secondary, downstream market: algorithmic auditing, bias detection, and compliance verification. If they cannot conduct their research without access to the proprietary weights, fine-tuning mechanisms, and high-volume APIs of the frontier models, the

LLM itself unequivocally becomes an essential facility.[25]
The arbitrary, unappealable throttling of a paid safety researcher therefore constitutes an illegal denial of an essential facility, provided it can be proven that the AI lab wields monopoly power and that providing access is technically feasible.[2] Given that the labs already provide commercial API access at scale, allocating un-throttled bandwidth to a paying researcher is highly feasible and poses no technical burden.[2] To combat these abuses, scholars and policymakers are increasingly advocating for an "AI Neutrality" mandate—akin to net neutrality rules for telecommunications—that would legally prohibit foundation model providers from unjustly discriminating, throttling, or artificially altering the latency and quality of service for third-party downstream users and auditors.[15]

Doctrine Element (MCI v. AT&T)	Application to AI Infrastructure & Safety Research
1. Monopolist Control	GAMMANs exert near-total, oligopolistic control over hyperscale compute and top-tier LLM API endpoints.
2. Inability to Duplicate	Multi-billion-dollar training costs, copyright data restrictions, and hardware shortages make replicating frontier models impossible for independent safety researchers.
3. Denial of Use	AI labs impose unappealable rolling limits, execute discriminatory API revocations (e.g., Windsurf), and utilize automated support walls to deny the high-volume access needed for rigorous auditing.
4. Feasibility of Provision	Labs already provide API access commercially; allocating un-throttled bandwidth to a paying researcher is highly technically feasible and incurs negligible marginal cost.

4. Regulatory Oversight: The FTC Act (Section 5) and Unfair Methods of Competition

Beyond the strict and heavily litigated strictures of the Sherman and Clayton Acts, federal regulators possess broader, more flexible authority under Section 5 of the Federal Trade Commission (FTC) Act. Section 5 explicitly prohibits "unfair methods of competition in or affecting commerce" alongside "unfair or deceptive acts or practices".[36] This statutory authority is deliberately designed to arrest anticompetitive behavior in its incipiency, granting the FTC the mandate to target conduct that violates the spirit of antitrust laws, even if it does not technically fulfill the rigid, judicially narrowed requirements of Section 1 or Section 2 of the Sherman Act.[22]

Current Agency Inquiries into Big Tech AI Monopolies

The FTC, under the leadership of Chair Lina M. Khan, has aggressively utilized its Section 5 enforcement authority and its Section 6(b) investigatory powers to scrutinize the highly concentrated AI landscape. Section 6(b) authorizes the Commission to conduct wide-ranging, mandatory studies that do not require a specific law enforcement purpose, compelling companies to submit detailed internal documents regarding their business practices.[37]

In early 2024, the FTC launched comprehensive 6(b) inquiries into the multi-billion-dollar investments and partnerships between the dominant cloud service providers (CSPs) and leading AI developers. Specifically, the orders targeted Microsoft's relationship with OpenAI, Amazon's partnership with Anthropic, and Google's partnership with Anthropic.[39] The resulting 2025 FTC staff report revealed profound concerns regarding these "co-opetition" dynamics.[1] The report highlighted that these partnerships often grant the CSPs extensive equity, revenue-sharing rights, and exclusive access to sensitive technical information.[40] This structure effectively creates "lock-in," deprives nascent startups of key AI inputs (such as engineering talent and compute resources), and reveals sensitive data that can systematically undermine fair competition.[40]

Through these complex investments, dominant tech firms create "walled gardens," utilizing restrictive Application Programming Interfaces (APIs) to exert total control over the downstream data flow and application layer.[22] The FTC has historically viewed the strategic closing of APIs to stifle dependent third parties as an unfair method of competition. Notable precedents include the FTC's Section 5 actions against Intel for redesigning software to purposefully disadvantage competing CPUs, and actions against Dell for reneging on information-sharing agreements essential for industry interoperability.[22] If major AI labs utilize their market dominance to arbitrarily throttle independent safety researchers, this behavior squarely aligns with the FTC's definition of an incipient tying arrangement or a breach of interoperability designed to "pick winners" and insulate the dominant platform from external scrutiny.[22]

Mechanisms for Filing Formal Complaints Regarding Algorithmic Throttling

For private citizens, independent labs, or safety researchers facing algorithmic throttling explicitly designed to suppress safety narratives, the federal government provides specific avenues for regulatory redress. While the FTC does not generally resolve individual private contractual disputes, it actively aggregates complaints to identify patterns of systemic anticompetitive behavior that warrant federal intervention and broad enforcement actions.[42]

An independent researcher experiencing unappealable API revocation, discriminatory throttling, or the imposition of exclusionary digital barriers can file a formal complaint directly with the **FTC's Bureau of Competition** via their dedicated antitrust complaint intake webform.[42] To maximize the efficacy of such a complaint, the researcher must frame the throttling not merely as a breach of terms of service or poor customer support, but

strategically as an "unfair method of competition" under Section 5 of the FTC Act. The complaint must meticulously document how the AI lab is leveraging its gatekeeping power to foreclose the secondary market of AI safety auditing. By doing so, the lab maintains proprietary control over public safety narratives and protects its core monopoly from regulatory or competitive disruption generated by negative audit findings.[19]

Concurrently, if the independent safety research is funded by government grants or involves federal procurement, complaints concerning anticompetitive impacts can be routed to the DOJ's Procurement Collusion Strike Force (PCSF).[44] Furthermore, researchers employed by these labs or affiliated entities who blow the whistle on these practices are protected from employer retaliation under the Criminal Antitrust Anti-Retaliation Act of 2019 (CAARA), and can file retaliation complaints directly with the Occupational Safety and Health Administration (OSHA).[45] By utilizing these formal mechanisms, researchers can trigger the FTC's immense investigative powers, including the issuance of Civil Investigative Demands (CIDs) and subpoenas under Section 9 of the FTC Act, to compel the disclosure of the proprietary algorithms used to enact the throttling.[37]

5. Discriminatory Monopolies: The Intersection of Antitrust Law and the ADA

The most novel, legally complex, and arguably insidious dimension of AI gatekeeping lies at the intersection of antitrust law and civil rights—specifically, Title III of the Americans with Disabilities Act (ADA) and Section 504 of the Rehabilitation Act. When a monopoly or a highly concentrated cartel possesses the absolute market power to dictate the technological standards of an entire industry, the adoption of exclusionary digital infrastructure transcends mere poor user experience. It becomes a systemic, structural barrier to entry that disproportionately and illegally affects protected classes, weaponizing digital exclusion as an anticompetitive tool.[47]

Digital Infrastructure Standards and Statutory Precedent

Title III of the ADA mandates that places of "public accommodation" ensure effective communication and equal access to goods and services for individuals with disabilities.[49] Historically, federal circuit courts have been split on whether online-only businesses lacking a physical brick-and-mortar storefront constitute public accommodations.[50] However, recent jurisprudence and sweeping regulatory updates by the DOJ have increasingly affirmed that digital platforms, APIs, and mobile applications must conform to strict accessibility standards, most notably the Web Content Accessibility Guidelines (WCAG) 2.1 Level AA.[49]

In April 2024, the DOJ published a final rule under Title II of the ADA mandating WCAG 2.1 AA compliance for state and local government digital properties.[52] Legal experts widely agree this Title II rule establishes the *de facto* standard that will soon be aggressively applied to private sector Title III entities.[47] Furthermore, in May 2024, the Department of Health and Human Services (HHS) finalized a historic rule under Section 504 of the Rehabilitation Act, explicitly requiring that all web content and mobile apps developed by organizations receiving federal

funding conform to WCAG 2.1 AA.[53] Section 504 explicitly treats inaccessible digital experiences as actionable disability discrimination.[55] Together, these precedents establish a clear federal mandate: digital infrastructure must be accessible. The failure to comply has resulted in a massive surge in litigation, with plaintiffs filing 4,187 digital accessibility lawsuits in 2024 alone.[48] Key cases, such as *Robles v. Domino's* and *Gil v. Winn-Dixie*, have cemented the legal reality that digital inaccessibility yields immense liability.[57]

Exclusionary Design as an Anticompetitive Moat

In the context of frontier AI, major laboratories often deploy highly automated, AI-driven support walls to handle API quota requests, account appeals, and throttling inquiries.[16] These automated systems are frequently designed without comprehensive accessibility overlays, utilizing complex visual CAPTCHAs, poorly coded dynamic web elements, or strictly mouse-dependent interfaces that are entirely incompatible with screen readers, voice navigation, or other assistive technologies relied upon by disabled researchers.[58]

When a competitive market functions properly, a disabled user facing an inaccessible interface can simply take their business to a competitor offering an accessible alternative. However, in a monopolistic or perfectly cartelized market, there are no viable alternatives.[26] If the GAMMANs universally adopt inaccessible AI support walls, they effectively establish an industry-wide standard that systematically excludes disabled populations from participating in the AI economy, without offering any alternatives.[26]

This digital exclusion serves a profound, deeply anticompetitive purpose. The sociological and legal theory regarding the "aesthetics of disability" posits that structural discrimination often relies on minimizing contact between disabled individuals and the dominant economic infrastructure.[60] By deploying inaccessible digital choke-points, AI monopolies functionally sanitize their user base, filtering out individuals who require customized, high-touch support or who might deploy specialized accessibility scripts that algorithms misidentify as "high-volume" or "bot" API usage.

When an independent safety researcher with a visual or motor disability is throttled by an AI lab's opaque "rolling limit" and subsequently trapped behind an inaccessible, automated appeal system, the injury is twofold. First, the researcher is denied the essential facility of the LLM, entirely halting their vital auditing work.[25] Second, and more critically from an antitrust perspective, the monopolistic firm insulates itself from the legal, competitive, and public relations pressures that a diverse, independent auditing body would inevitably generate. By passively enforcing a digital standard that violates WCAG 2.1 AA and ADA Title III, the AI cartel homogenizes the pool of actors capable of interacting with its systems, locking out disabled researchers who might uncover novel algorithmic biases or safety vulnerabilities.[48]

This intersection creates a compelling, multi-front legal vulnerability for Big Tech. Antitrust regulators (DOJ/FTC) and civil rights litigators can converge on the premise that the failure to provide accessible alternatives is not merely a technical oversight or a basic ADA violation, but a calculated, exclusionary standard-setting practice. When this failure is engineered by a monopoly to lock out a specific subset of researchers, it bridges the gap between civil rights violations and exclusionary anticompetitive conduct under Section 2 of the Sherman Act,

demonstrating how market power can be weaponized to evade both accessibility mandates and independent safety oversight.

Legal Framework	Regulatory Focus	Implication for AI Gatekeeping	Mechanism of Exclusion
ADA Title III	Public accommodations must provide equal digital access (WCAG 2.1 AA).	AI support walls must be fully navigable by assistive tech.	Visual CAPTCHAs and mouse-only interfaces trap disabled researchers.
Rehabilitation Act (Sec. 504)	Prohibits discrimination by entities receiving federal funds.	Federally backed AI research grants require accessible API interfaces.	Failure to provide accessible endpoints violates federal funding terms.
Sherman Act (Sec. 2)	Prohibits the maintenance of monopoly power via exclusionary conduct.	Inaccessible standards act as an anticompetitive moat.	Monopolists use non-compliant design to filter out auditors, insulating the model from scrutiny.

6. Comprehensive Synthesis and Strategic Outlook

The rapid consolidation of the artificial intelligence sector, driven by insurmountable capital requirements and massive compute bottlenecks, demands a robust and multifaceted regulatory response. As this exhaustive analysis demonstrates, the collective behavior of major AI laboratories—specifically the arbitrary throttling of API access, the deployment of unappealable automated support walls, and the leveraging of exclusive hyperscale computational advantages—violates the foundational principles of competition law across multiple statutory dimensions.

1. **Collusion and Coordination:** Under Section 1 of the Sherman Act, the uniform, industry-wide adoption of restrictive API limits and automated blockades can be aggressively prosecuted as a hub-and-spoke conspiracy. This requires demonstrating that labs rely on shared algorithmic infrastructure or act against their independent economic interests to maintain a collective boycott of independent safety researchers, thereby chilling the secondary market for alignment auditing.[3]
2. **Monopolization of Compute:** Under Section 2 of the Sherman Act, the absolute control over hyperscale compute and foundational models by a highly concentrated oligopoly constitutes an actionable shared monopoly under modern frameworks like the Moligopoly thesis.[1] The weaponization of this power to choke compute supply and degrade downstream applications—as explicitly alleged in the recent *Bryant v. Microsoft* litigation—illustrates a clear pattern of exclusionary conduct intended to neutralize competitive threats and suppress oversight.[28]
3. **Essential Facilities:** Because the multi-billion-dollar cost of training frontier models renders duplication practically impossible, LLMs and their underlying cloud infrastructure

act as modern essential facilities. The discriminatory revocation of access, exemplified by the 2025 Anthropic/Windsurf incident, establishes a dangerous precedent of dominant firms starving downstream competitors and safety auditors of essential digital public utilities.[2]

4. **FTC Enforcement:** The FTC's broad authority under Section 5 of the FTC Act perfectly encapsulates the incipient threat of AI "walled gardens" and the lock-in effects of CSP partnerships. Formal complaints to the Bureau of Competition regarding unfair algorithmic throttling offer a direct mechanism for researchers to challenge the labs' efforts to exercise proprietary control over the public safety narrative.[22]
5. **Civil Rights Intersection:** The most profound regulatory insight emerges from the synthesis of antitrust principles and the ADA. By embedding inaccessible digital standards into the very support walls required to appeal API throttling, monopolies systematically exclude disabled researchers.[48] This conduct transcends traditional discrimination; it functions as a calculated anticompetitive moat that actively insulates dominant models from rigorous, diverse, and independent algorithmic auditing, shielding the gatekeepers from accountability.

To preserve the integrity of the technological frontier and ensure that the trajectory of General-Purpose AI is subjected to independent ethical alignment, regulatory bodies must look beyond traditional price-theory antitrust enforcement. They must embrace the deeply interconnected nature of digital civil rights, essential facilities mandates, and algorithmic collusion to systematically dismantle the gatekeeping mechanisms currently insulating the world's most powerful AI laboratories.

Works cited

1. Why artificial intelligence is creating fundamental challenges for competition policy - Bruegel, accessed April 16, 2026, https://www.bruegel.org/policy-brief/why-artificial-intelligence-creating-fundamental-challenges-competition-policy
2. An Antimonopoly Approach to Governing Artificial Intelligence - Yale Law & Policy Review, accessed April 16, 2026, https://yalelawandpolicy.org/antimonopoly-approach-governing-artificial-intelligence
3. AI and Antitrust: “The Algorithm Made Me Do It”, accessed April 16, 2026, https://competition.scholasticahq.com/article/124368-ai-and-antitrust-the-algorithm-made-me-do-it
4. Antitrust 101: A Quick Spin Through Hub-and-Spoke Conspiracies | Winston & Strawn, accessed April 16, 2026, https://www.winston.com/en/blogs-and-podcasts/competition-corner/antitrust-101-a-quick-spin-through-hub-and-spoke-conspiracies
5. Hub-and-Spoke Conspiracies: Can Big Data and Pricing Algorithms Form the Rim? - SMU Scholar, accessed April 16, 2026, https://scholar.smu.edu/cgi/viewcontent.cgi?article=1350&context=scitech
6. Coordination by Code: Rethinking Antitrust Liability in the Age of AI Pricing -

American Bar Association, accessed April 16, 2026, https://www.americanbar.org/groups/antitrust_law/resources/magazine/2025-fall/coordination-code-rethinking-antitrust-liability-age-ai-pricing/
7. Exhibit 1 Comments of Center for Democracy & Technology: U.S. and Plaintiff States v. RealPage, Inc. - Department of Justice, accessed April 16, 2026, https://www.justice.gov/atr/media/1398941/dl
8. SUSTAINABLE AND UNCHALLENGED ALGORITHMIC TACIT COLLUSION - Scholarly Commons: Northwestern Pritzker School of Law, accessed April 16, 2026, https://scholarlycommons.law.northwestern.edu/cgi/viewcontent.cgi?article=1331&context=njtip
9. Tacit Agreements to Collude: Enforcing Section 1 of the Sherman Act in the Age of Algorithms - American Bar Association, accessed April 16, 2026, https://www.americanbar.org/groups/antitrust_law/resources/magazine/2025-fall/tacit-agreements-collude-enforcing-section-1-sherman-act-age-algorithms/
10. Algorithmic Price-Fixing: US States Hit Control-Alt-Delete on Digital Collusion | Perkins Coie, accessed April 16, 2026, https://perkinscoie.com/insights/update/algorithmic-price-fixing-us-states-hit-control-alt-delete-digital-collusion
11. Clarity Starts to Form Regarding Pricing Algorithm Antitrust Cases - Steptoe, accessed April 16, 2026, https://www.steptoe.com/en/news-publications/stepahead-antitrust-and-competition-insights/clarity-starts-to-form-regarding-pricing-algorithm-antitrust-cases.html
12. Artificial Intelligence and Antitrust: When Do Algorithms Violate Competition Laws? I. What's the issue? - Quinn Emanuel, accessed April 16, 2026, https://www.quinnemanuel.com/media/vdwbb1ag/client-alert-artificial-intelligence-and-antitrust-when-do-algorithms-violate-competition-laws.pdf
13. Can Machines Commit Crimes Under U.S. Antitrust Laws?, accessed April 16, 2026, https://businesslawreview.uchicago.edu/print-archive/can-machines-commit-crimes-under-us-antitrust-laws
14. Artificial Intelligence and Antitrust: When Do Algorithms Violate Competition Laws?, accessed April 16, 2026, https://www.quinnemanuel.com/the-firm/publications/artificial-intelligence-and-antitrust-when-do-algorithms-violate-competition-laws/
15. AI Neutrality - Vanderbilt University, accessed April 16, 2026, https://cdn.vanderbilt.edu/vu-URL/wp-content/uploads/sites/412/2026/01/22160718/AI-Neutrality.pdf
16. Tell HN: Twitter switched temporarily to rate limited mode | Hacker News, accessed April 16, 2026, https://news.ycombinator.com/item?id=36552324
17. International Metals Reclamation Company (INMETCO) RACT 2 SIP Summary TV-37-00243 Ellwood City Borough, Lawrence County Page Nu - Regulations.gov, accessed April 16, 2026, https://downloads.regulations.gov/EPA-R03-OAR-2021-0380-0002/attachment_7

.pdf
18. Tell HN: Anthropic no longer allowing Claude Code subscriptions to use OpenClaw | Hacker News, accessed April 16, 2026, https://news.ycombinator.com/item?id=47633396
19. How Antitrust Can Promote AI Safety Collaborations - Lawfare, accessed April 16, 2026, https://www.lawfaremedia.org/article/how-antitrust-can-promote-ai-safety-collaborations
20. Breiding v. Eversource Energy & Avangrid, Inc. | Civil Action No. 17-12274 | D. Mass. | Judgment | Law | CaseMine, accessed April 16, 2026, https://www.casemine.com/judgement/us/5b99dbeb342cca319d1fc32d
21. The Antitrust Case Against AI Overviews - Harvard Journal of Law & Technology, accessed April 16, 2026, https://jolt.law.harvard.edu/digest/the-antitrust-case-against-ai-overviews
22. Concentrated Digital Markets, Restrictive APIs, and the Fight for Internet Interoperability - The University of Memphis, accessed April 16, 2026, https://www.memphis.edu/law/documents/04_sharma.pdf
23. Antitrust Analysis: Problems, Text, and Cases, 8th Edition, 2021, by Areeda, Kaplow, Edlin, and Hemphill [7 ed.] 9781454824992 - DOKUMEN.PUB, accessed April 16, 2026, https://dokumen.pub/antitrust-analysis-problems-text-and-cases-8th-edition-2021-by-areeda-kaplow-edlin-and-hemphill-7nbsped-9781454824992.html
24. Cornpetition Policy for S m d Market Economies: Market Conditions Under the Magnifying Glass, accessed April 16, 2026, https://www.collectionscanada.ca/obj/s4/f2/dsk2/ftp03/NQ53773.pdf
25. General-Purpose AI Models as Essential Inputs in Downstream Markets: The Need for a Strict Standard Regarding Mandatory Access | GRUR International | Oxford Academic, accessed April 16, 2026, https://academic.oup.com/grurint/article/73/10/948/7742781
26. U.S. v. Google — A Landmark Case | Purdue Global Law School, accessed April 16, 2026, https://www.purduegloballawschool.edu/blog/news/google-landmark-case
27. In Landmark Decision, D.C. Federal Court Holds Google Maintained an Illegal Monopoly in Internet Search and Advertising Markets and Sets the Stage For Future Enforcement Actions | White & Case LLP, accessed April 16, 2026, https://www.whitecase.com/insight-our-thinking/landmark-decision-dc-federal-court-holds-google-maintained-illegal-monopoly
28. 1 2 3 4 5 6 7 8 9 10 11 12 13 14 15 16 17 18 19 20 21 ... - Class Action, accessed April 16, 2026, https://www.classaction.org/media/bryant-et-al-v-microsoft-corporation-complaint.pdf
29. THE SEPARATION OF PLATFORMS AND COMMERCE - Columbia Law Review -, accessed April 16, 2026, https://columbialawreview.org/content/the-separation-of-platforms-and-commerce/
30. Nikolas Guggenberger* Digital platforms are the railroads of the modern era. In

the early twentieth century, a vast railroad net, accessed April 16, 2026, https://appliedantitrust.com/16_foreclosure/digital_markets/general/guggenberger_essential_platforms2020.pdf
31. Liability for Data Scraping Prohibitions under the Refusal to Deal Doctrine: An Incremental Step toward More Robust Sherman Act Enforcement | The University of Chicago Law Review, accessed April 16, 2026, https://lawreview.uchicago.edu/print-archive/liability-data-scraping-prohibitions-under-refusal-deal-doctrine-incremental-step
32. The Essential Facilities Doctrine in the Digital Economy: Dispelling Persistent Myths, accessed April 16, 2026, https://yjolt.org/essential-facilities-doctrine-digital-economy-dispelling-persistent-myths
33. Essential Platforms | Stanford Law School, accessed April 16, 2026, https://law.stanford.edu/wp-content/uploads/2021/05/publish_this_-_guggenberger_essential_platforms_eic.pdf
34. Android Auto: the end of the essential facility doctrine as we know it - Wolters Kluwer, accessed April 16, 2026, https://legalblogs.wolterskluwer.com/competition-blog/android-auto-the-end-of-the-essential-facility-doctrine-as-we-know-it/
35. CLOUD INFRASTRUCTURE-AS-A-SERVICE AS AN ESSENTIAL FACILITY: MARKET STRUCTURE, COMPETITION, AND THE NEED FOR INDUSTRY AND REGULAT - Berkeley Technology Law Journal, accessed April 16, 2026, https://btlj.org/data/articles2019/34_1/03_Benzina_Web.pdf
36. Antitrust Year in Review and a Look Ahead to 2025 - Hogan Lovells, accessed April 16, 2026, https://www.hoganlovells.com/en/publications/antitrust-yearinreview-and-look-ahead-to-2025
37. A Brief Overview of the Federal Trade Commission's Investigative, Law Enforcement, and Rulemaking Authority, accessed April 16, 2026, https://www.ftc.gov/about-ftc/mission/enforcement-authority
38. Unfairness, Reconstructed - Yale Law School Open Scholarship Repository, accessed April 16, 2026, https://openyls.law.yale.edu/server/api/core/bitstreams/547850f9-0d87-4cbe-b1c2-e2b7bd1dc0b3/content
39. FTC Launches Inquiry into Generative AI Investments and Partnerships, accessed April 16, 2026, https://www.ftc.gov/news-events/news/press-releases/2024/01/ftc-launches-inquiry-generative-ai-investments-partnerships
40. FTC Issues Staff Report on AI Partnerships & Investments Study, accessed April 16, 2026, https://www.ftc.gov/news-events/news/press-releases/2025/01/ftc-issues-staff-report-ai-partnerships-investments-study
41. Partnerships Between Cloud Service Providers and AI Developers - Federal Trade Commission, accessed April 16, 2026, https://www.ftc.gov/system/files/ftc_gov/pdf/p246201_aipartnerships6breport_re

dacted_0.pdf
42. Antitrust Complaint Intake - Federal Trade Commission, accessed April 16, 2026, https://www.ftc.gov/advice-guidance/competition-guidance/antitrust-complaint-intake
43. Contact the Federal Trade Commission, accessed April 16, 2026, https://www.ftc.gov/about-ftc/contact
44. PCSF Citizen Complaint - Antitrust Division - Department of Justice, accessed April 16, 2026, https://www.justice.gov/atr/webform/pcsf-citizen-complaint
45. Submit Your Antitrust Report Online | United States Department of Justice, accessed April 16, 2026, https://www.justice.gov/atr/webform/submit-your-antitrust-report-online
46. Report Antitrust Concerns to the Antitrust Division | United States Department of Justice, accessed April 16, 2026, https://www.justice.gov/atr/report-antitrust-concerns
47. DOJ Final Rule on Website Accessibility for State and Local Governments Portends Significant Changes for Private-Sector Websites - Ogletree Deakins, accessed April 16, 2026, https://ogletree.com/insights-resources/blog-posts/doj-final-rule-on-website-accessibility-for-state-and-local-governments-portends-significant-changes-for-private-sector-websites/
48. ADA Website Lawsuits 2026: Trends & Defense - Kha Creation, accessed April 16, 2026, https://khacreationusa.com/ada-title-iii-website-lawsuits/
49. Digital Accessibility Under Title III of the ADA: Recent Developments and Risk Mitigation Best Practices - American Bar Association, accessed April 16, 2026, https://www.americanbar.org/groups/business_law/resources/business-law-today/2025-august/digital-accessibility-under-title-iii-ada/
50. Minnesota District Court Says Web-Only Businesses Are Subject to Title III of the ADA, accessed April 16, 2026, https://www.adatitleiii.com/2025/03/minnesota-district-court-says-web-only-businesses-are-subject-to-title-iii-of-the-ada/
51. ADA Website Compliance Lawsuits | 2024 High Profile Cases - Level Access, accessed April 16, 2026, https://www.levelaccess.com/blog/title-iii-lawsuits-10-big-companies-sued-over-website-accessibility/
52. Fact Sheet: New Rule on the Accessibility of Web Content and Mobile Apps Provided by State and Local Governments | ADA.gov, accessed April 16, 2026, https://www.ada.gov/resources/2024-03-08-web-rule/
53. HHS Section 504 Digital Accessibility Rule Is in Effect. What To Do If You're Not Compliant, accessed April 16, 2026, https://testpros.com/compliance/hhs-section-504-digital-accessibility-rule/
54. HHS's section 504 final rule includes digital accessibility standards - Reed Smith LLP, accessed April 16, 2026, https://www.reedsmith.com/articles/hhss-section-504-final-rule-includes-digital-accessibility-standards/
55. What Happens When Section 504 Digital Accessibility Requirements are Violated,

accessed April 16, 2026, https://www.siteimprove.com/blog/section-504-digital-accessibility-requirements-violations/
56. A Rise in ADA Website Accessibility Lawsuits May Leave You Asking: Is My Website A Risk?, accessed April 16, 2026, https://www.saul.com/insights/blog/ada-website-accessibility-risk
57. ADA Compliance Requirements for Websites 2026 | A Must-Read Guide - DigitalA11Y, accessed April 16, 2026, https://www.digitala11y.com/compliance/ada/
58. ADA Title III | Disability Law for Businesses | Seyfarth Shaw, accessed April 16, 2026, https://www.adatitleiii.com/
59. Research Brief: Digital Access and Title III of the ADA, accessed April 16, 2026, https://adata.org/research_brief/digital-access-and-title-iii-ada
60. THE AESTHETICS OF DISABILITY - Columbia Law Review, accessed April 16, 2026, https://www.columbialawreview.org/content/the-aesthetics-of-disability/